

Assignment 15 sep

Perform Data preprocessing on Titanic dataset 1.Data Collection. Please download the dataset from

<https://www.kaggle.com/datasets/yasserh/titanic-dataset>

2.Data Preprocessing o Import the Libraries. o Importing the dataset. o Checking for Null Values. o Data Visualization. o Outlier Detection o Splitting Dependent and Independent variables o Perform Encoding o Feature Scaling. o Splitting Data into Train and Test

1.Data Collection :

Data Set is collected from the kaggle website

2.Data Preprocessing :

Importing the Libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Importing the DataSet

```
df=pd.read_csv("Titanic-Dataset.csv")
df.head()
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs T. B.)	female	38.0	1	0	PC 17599

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   PassengerId  891 non-null    int64
1   Survived     891 non-null    int64
2   Pclass       891 non-null    int64
3   Name         891 non-null    object
4   Sex          891 non-null    object
5   Age          714 non-null    float64
6   SibSp        891 non-null    int64
7   Parch        891 non-null    int64
8   Ticket       891 non-null    object
9   Fare         891 non-null    float64
10  Cabin        204 non-null    object
11  Embarked     889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
df.describe()
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	
count	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.000000
std	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000

Checking for Null Values

50%	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.400000
-----	------------	----------	----------	-----------	----------	----------	-----------

df.isnull().any()

PassengerId	False
Survived	False
Pclass	False
Name	False
Sex	False
Age	True
SibSp	False
Parch	False
Ticket	False
Fare	False
Cabin	True
Embarked	True
dtype: bool	

df.isnull().sum()

PassengerId	0
Survived	0
Pclass	0
Name	0
Sex	0
Age	177
SibSp	0
Parch	0
Ticket	0
Fare	0
Cabin	687
Embarked	2
dtype: int64	

```
print("Null percentage in columns : ")
for i in df.columns:
    c=df[i].count()
    n=df[i].isnull().sum()
    print(i," : ",(n/(n+c)) * 100)
```

Null percentage in columns :	
PassengerId	: 0.0
Survived	: 0.0
Pclass	: 0.0
Name	: 0.0
Sex	: 0.0
Age	: 19.865319865319865
SibSp	: 0.0
Parch	: 0.0
Ticket	: 0.0
Fare	: 0.0
Cabin	: 77.10437710437711
Embarked	: 0.22446689113355783

df.shape

(891, 12)

```
df["Age"].fillna(df["Age"].median(),inplace=True)
df["Embarked"].fillna(df["Embarked"].mode()[0],inplace=True)
```

```
print(df["Age"].isnull().any())
print(df["Embarked"].isnull().any())
```

False
False

df.head()

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171
				Cumings, Mrs. John					

```
print(df.shape)

(891, 12)
```

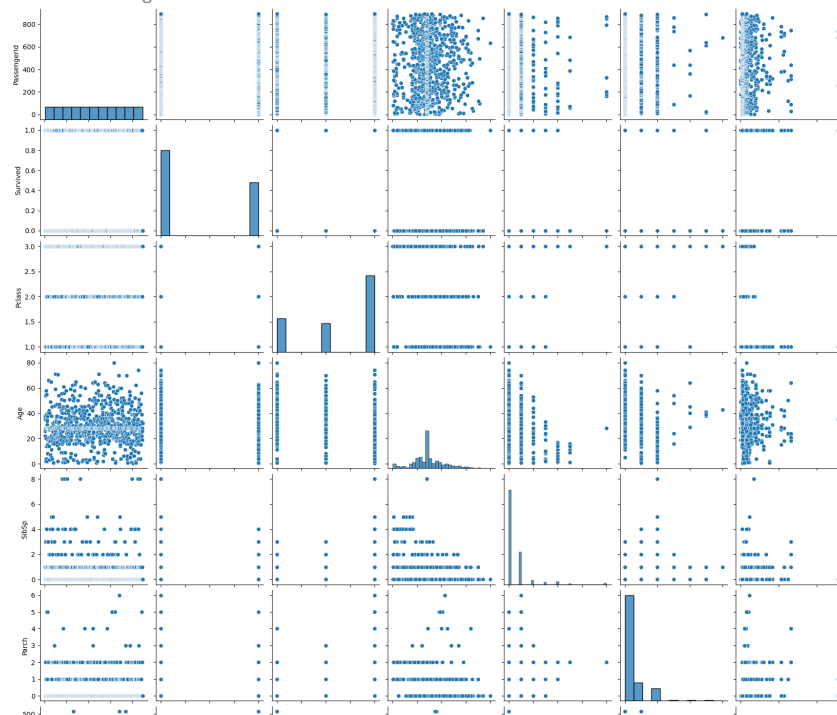
```
df.isnull().any()

PassengerId    False
Survived        False
Pclass          False
Name            False
Sex             False
Age            False
SibSp           False
Parch           False
Ticket          False
Fare            False
Cabin           True
Embarked        False
dtype: bool
```

▼ Data Visualization


```
sns.pairplot(df)
```

```
<seaborn.axisgrid.PairGrid at 0x7c24a402e9e0>
```



```
sns.scatterplot(x="Embarked", y="Age", data=df)
```

```
<Axes: xlabel='Embarked', ylabel='Age'>
```




```
sns.barplot(x="Sex", y="Survived", data=df, hue="Pclass")
```

```
<Axes: xlabel='Sex', ylabel='Survived'>
```


```
sns.barplot(x="Sex",y="Survived",data=df,hue="Parch")
```

```
<Axes: xlabel='Sex', ylabel='Survived'>
```




```
sns.barplot(x="Sex",y="SibSp",data=df,hue="Pclass")
```

```
<Axes: xlabel='Sex', ylabel='SibSp'>
```



```
sns.barplot(x="Embarked",y="Survived",data=df)
```

```
<Axes: xlabel='Embarked', ylabel='Survived'>
```



```
sns.distplot(df["Survived"])
```

<ipython-input-22-6525837c6049>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df["Survived"])
<Axes: xlabel='Survived', ylabel='Density'>
```



```
corr=df.corr(numeric_only=True)
```

```
corr
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	F
PassengerId	1.000000	-0.005007	-0.035144	0.034212	-0.057527	-0.001652	0.012
Survived	-0.005007	1.000000	-0.338481	-0.064910	-0.035322	0.081629	0.257
Pclass	-0.035144	-0.338481	1.000000	-0.339898	0.083081	0.018443	-0.549
Age	0.034212	-0.064910	-0.339898	1.000000	-0.233296	-0.172482	0.096
SibSp	-0.057527	-0.035322	0.083081	-0.233296	1.000000	0.414838	0.159
Parch	-0.001652	0.081629	0.018443	-0.172482	0.414838	1.000000	0.216
Fare	0.012658	0.257307	-0.549500	0.096688	0.159651	0.216225	1.000


```
sns.heatmap(corr,annot=True)
```




Outlier Detection



```
sns.boxplot(df)
```



```
sns.boxplot(data=df, x="Age", y="Embarked")
```




```
sns.boxplot(data=df, x="Fare", y="Embarked")
```

<Axes: xlabel='Fare', ylabel='Embarked'>



```
sns.boxplot(df["Fare"])
```

<Axes: >



```
df["Age"].skew()
```

```
0.5102446555756495
```

```
df["Fare"].skew() # as skewness should be -1 to +1 is normal range but here we are having so much outliers and should be treated first
```

```
4.787316519674893
```

```
df["Fare"].median()
```

```
14.4542
```

```
Q1 = df['Fare'].quantile(0.25)
```

```
Q3 = df['Fare'].quantile(0.75)
```

```
IQR = Q3 - Q1
```

```
width = 1.5
```


```
lower_limit = Q1 - (width*IQR)
```

```
upper_limit = Q3 + (width*IQR)
```

```
df['Fare'] = np.where(df['Fare'] > upper_limit, 14.4542, np.where(df['Fare'] < lower_limit, 14.4542, df["Fare"]))
```


```
sns.boxplot(df["Fare"])
```

<Axes: >



```
sns.boxplot(df.Age)
```


<Axes: >




```
print(df.Age.median())  
print(df.Age.shape)
```

```
28.0  
(891,)
```

```
p=df["Age"].quantile(0.99)  
p1=df['Age'].quantile(0.01)  
df=df[df['Age']<=p]  
df=df[df['Age']>=p1]
```

```
sns.boxplot(df.Age)
```

<Axes: >



```
sns.boxplot(df)
```



df.head()

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs	female	38.0	1	0	PC 17599

df.shape

(876, 12)

Splitting Dependent and Independent variables

```
df.drop(["PassengerId", "Name", "Ticket", "Cabin"],axis=1,inplace=True)
df.head()
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	0	3	male	22.0	1	0	7.2500	S
1	1	1	female	38.0	1	0	14.4542	C
2	1	3	female	26.0	0	0	7.9250	S
3	1	1	female	35.0	1	0	53.1000	S
4	0	3	male	35.0	0	0	8.0500	S

```
X=df.iloc[:,1:]
y=df.iloc[:,1]
```

X.head()

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	3	male	22.0	1	0	7.2500	S
1	1	female	38.0	1	0	14.4542	C
2	3	female	26.0	0	0	7.9250	S
3	1	female	35.0	1	0	53.1000	S
4	3	male	35.0	0	0	8.0500	S

y.head()

	Survived
0	0
1	1
2	1
3	1
4	0

```
y=y.squeeze()
```

```
type(X)
```

pandas.core.frame.DataFrame

```
type(y)
```

```
pandas.core.series.Series
```

```
y.head()

0    0
1    1
2    1
3    1
4    0
Name: Survived, dtype: int64
```

▼ Perform Encoding

```
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
X["Sex"]=le.fit_transform(X["Sex"])
mapping1=dict(zip(le.classes_,range(len(le.classes_))))
X["Embarked"]=le.fit_transform(X["Embarked"])
mapping2=dict(zip(le.classes_,range(len(le.classes_))))

print("For Sex Column :",mapping1)
print("For Embarked Column :",mapping2)

For Sex Column : {'female': 0, 'male': 1}
For Embarked Column : {'C': 0, 'Q': 1, 'S': 2}
```

X

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	3	1	22.0	1	0	7.2500	2
1	1	0	38.0	1	0	14.4542	0
2	3	0	26.0	0	0	7.9250	2
3	1	0	35.0	1	0	53.1000	2
4	3	1	35.0	0	0	8.0500	2
...
886	2	1	27.0	0	0	13.0000	2
887	1	0	19.0	0	0	30.0000	2
888	3	0	28.0	1	2	23.4500	2
889	1	1	26.0	0	0	30.0000	0
890	3	1	32.0	0	0	7.7500	1

876 rows × 7 columns

```
y

0    0
1    1
2    1
3    1
4    0
..
886   0
887   1
888   0
889   1
890   0
Name: Survived, Length: 876, dtype: int64
```

▼ Feature Scaling

```
from sklearn.preprocessing import StandardScaler
ss=StandardScaler()
X_Scale=pd.DataFrame(ss.fit_transform(X),columns=X.columns)

X_Scale.head()
```



	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	0.821711	0.743768	-0.589744	0.430836	-0.467137	-0.793895	0.582594
1	-1.573693	-1.344504	0.719906	0.430836	-0.467137	-0.227705	-1.956472
2	0.821711	-1.344504	-0.262332	-0.472065	-0.467137	-0.740846	0.582594
3	-1.573693	-1.344504	0.474346	0.430836	-0.467137	2.809531	0.582594

```
y.head()

0    0
1    1
2    1
3    1
4    0
Name: Survived, dtype: int64
```

```
X_Scale.shape

(876, 7)
```

```
y.shape

(876,)
```

Splitting Data into Train and Test

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X_Scale,y,test_size=0.2,random_state=0)
```

```
print(X_train,"\n",X_test,"\n","\n",y_train,"\n",y_test)

      Pclass      Sex      Age      SibSp      Parch      Fare      Embarked
45    0.821711    0.743768 -0.098626    0.430836 -0.467137 -0.145514 -0.686939
172   0.821711    0.743768 -0.098626    2.236638    0.772900    0.637785    0.582594
492   0.821711    0.743768 -0.426038 -0.472065 -0.467137 -0.751000    0.582594
820   0.821711    0.743768 -0.917157 -0.472065 -0.467137 -0.711374    0.582594
651  -1.573693    0.743768    1.702144    1.333737 -0.467137 -0.227705    0.582594
..      ...      ...      ...      ...      ...      ...      ...
835  -1.573693 -1.344504 -0.098626    0.430836 -0.467137 -0.227705 -1.956472
192   0.821711    0.743768 -0.098626 -0.472065 -0.467137 -0.754599 -0.686939
629  -0.375991    0.743768    0.146934    0.430836    0.772900    0.699346    0.582594
559   0.821711    0.743768 -0.835304 -0.472065 -0.467137 -0.743141    0.582594
684  -1.573693    0.743768    2.520675 -0.472065 -0.467137    0.722923    0.582594

[700 rows x 7 columns]
      Pclass      Sex      Age      SibSp      Parch      Fare      Embarked
141  -0.375991    0.743768 -0.835304    0.430836    0.772900    1.524558    0.582594
113  -0.375991    0.743768 -0.016772    0.430836 -0.467137    0.286740    0.582594
730  -1.573693    0.743768 -0.098626 -0.472065 -0.467137    0.994065    0.582594
294  -1.573693    0.743768 -0.098626 -0.472065 -0.467137    1.033360    0.582594
261  -0.375991    0.743768    0.556200 -0.472065 -0.467137 -0.538472    0.582594
..      ...      ...      ...      ...      ...      ...      ...
578  -1.573693 -1.344504 -0.917157 -0.472065    2.012937 -0.227705    0.582594
773   0.821711    0.743768 -0.344185 -0.472065 -0.467137 -0.793895    0.582594
522  -0.375991    0.743768 -0.507891    1.333737    0.772900 -0.459881    0.582594
780   0.821711 -1.344504 -0.098626    6.751142    2.012937 -0.227705    0.582594
54   -1.573693    0.743768 -0.098626 -0.472065 -0.467137    1.426319    0.582594

[176 rows x 7 columns]

46    0
176   0
499   0
834   0
660   1
..
849   1
196   0
637   0
566   0
694   0
Name: Survived, Length: 700, dtype: int64
145   0
117   0
740   1
298   1
265   0
..
585   1
785   0
```

```
529    0
792    0
55     1
Name: Survived, Length: 176, dtype: int64
```

```
print(X_train.shape,X_test.shape,y_train.shape,y_test.shape)
```

```
(700, 7) (176, 7) (700,) (176,)
```

▼ Preprocessing Done

▼ Testing for accuracy

```
from sklearn.linear_model import LogisticRegression
lr=LogisticRegression()
lr.fit(X_train,y_train)
```

```
▼ LogisticRegression
LogisticRegression()
```

```
y_predict=lr.predict(X_test)
```

```
from sklearn.metrics import accuracy_score
accuracy = accuracy_score(y_test, y_predict)
```

```
print("Accuracy:", accuracy)
```

```
Accuracy: 0.8295454545454546
```